



Programming Models Case Study: OpenMP

Ruini Xue


School of Computer Science and Engineering

University of Electronic Science and Technology of China

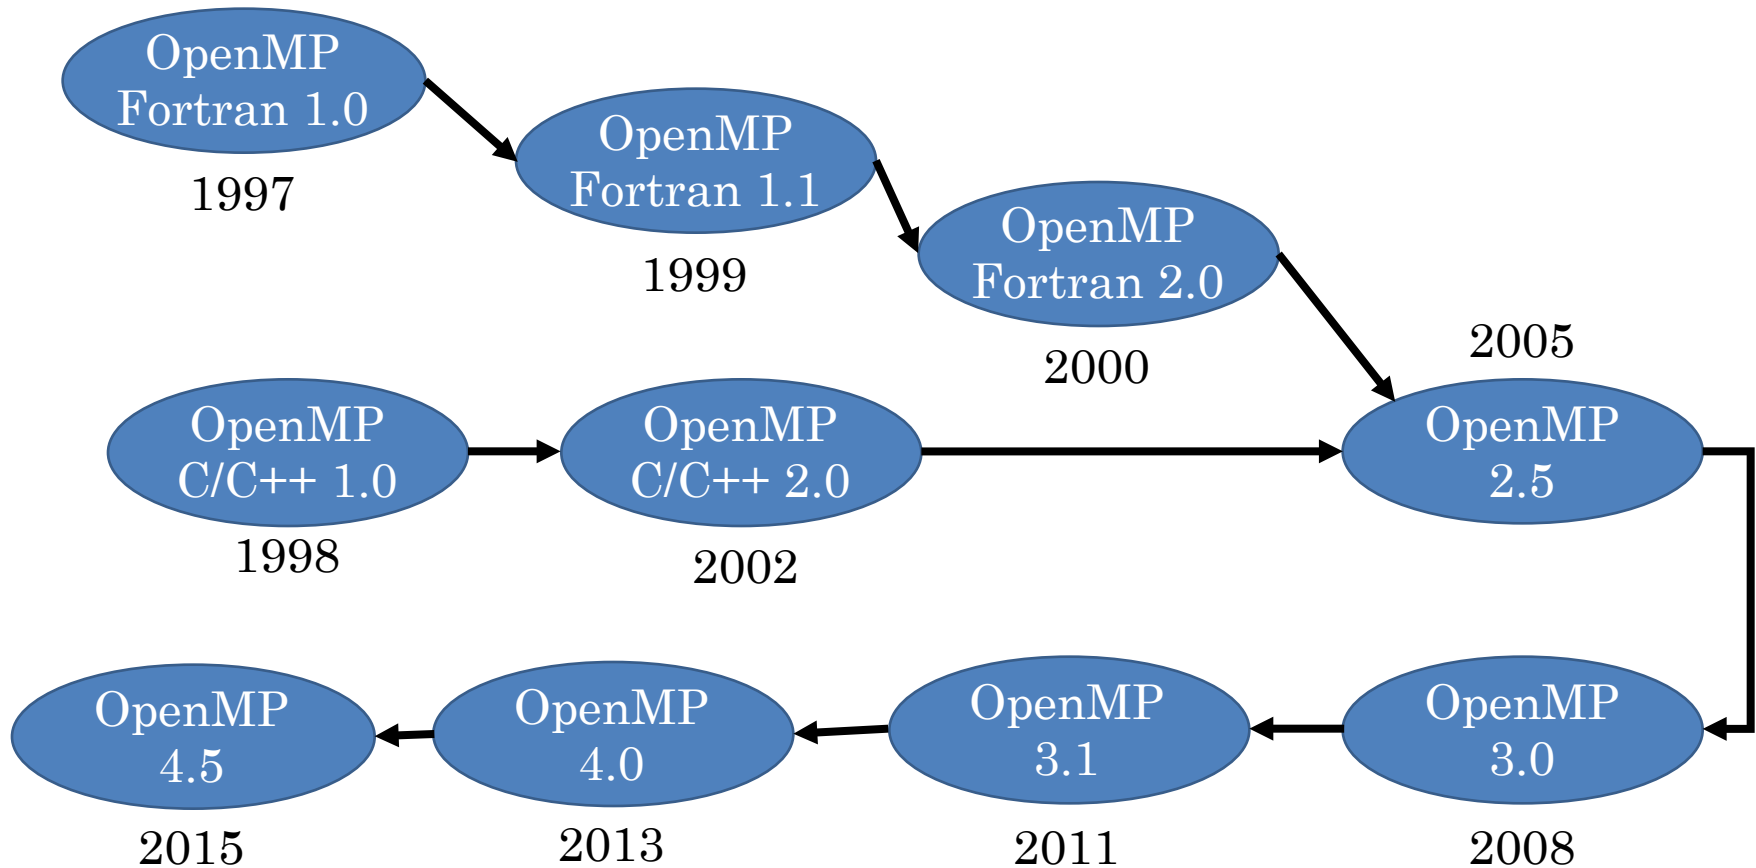
2016

What is OpenMP?

- OpenMP is an API for parallel programming
- Designed for shared-memory multiprocessors
- Set of compiler directives, library functions, and environment variables, but not a language
- Can be used with C, C++, or Fortran
- Based on fork/join model of threads



Release History



Fork/Join Programming Model



- When program begins execution, only master thread is active
- Master thread executes sequential portions of the program
- For parallel portions of program, master thread forks (creates or awakens) additional threads
- At join (end of parallel section of code), extra threads are suspended or die

Relating Fork/Join to Code



for {




}



for {



}



Sequential code

Parallel code


Sequential code

Parallel code

Sequential code

Another view

- All OpenMP programs begin as a single process: the master thread.
 - The master thread executes sequentially until the first parallel region construct is encountered
- Master thread spawns a team of threads as needed
- Parallelism added incrementally until performance goals are met: i.e. the sequential program evolves into a parallel program



Incremental Parallelization



- Sequential program is a special case of threaded program
- Programmers can add parallelism incrementally
- Profile program execution
- Repeat
 - Choose best opportunity for parallelization
 - Transform sequential code into parallel code
- Until further improvements not worth the effort
- Difficult for distributed memory programming

#pragma



- Directive is the method specified by the C standard for providing additional information to the compiler
 - it tells the compiler to do something, set some option, take some action, override some default, etc.
- E.g.

```
#pragma GCC error "message"
```

```
#pragma once  
// header file code
```

```
#pragma warning (disable : 4018 )
```


Syntax of Compiler Directives



- A C/C++ compiler directive is called a *pragma*
- Pragmas are handled by the preprocessor
- All OpenMP pragmas have the syntax:
 - `#pragma omp <rest of pragma>`
 - *structured block*
- Pragmas appear immediately before relevant construct

Pragma: parallel for

- tells the compiler that the for loop which immediately follows can be executed in parallel
- The number of loop iterations must be computable at run time before loop executes
- Loop must not contain a `break`, `return`, or `exit`
- Loop must not contain a `goto` to a label outside loop

parallel for [2.11.1] [2.10.1]

Shortcut for specifying a **parallel** construct containing one or more associated loops and no other statements.

```
#pragma omp parallel for [clause[ , clause] ...]  
    for-loop
```

Example



- `int first, *marked, prime, size;`
- `...`
- `#pragma omp parallel for`
- `for (i = first; i < size; i += prime)`
- `marked[i] = 1;`

1. Threads are assigned an independent set of iterations
2. Barrier/join: threads must wait at the end of construct

Pragma: parallel

- Sometimes the code that should be executed in parallel goes beyond a single `for` loop
- The `parallel` pragma is used when a block of code should be executed in parallel

parallel [2.5] [2.5]

Forms a team of threads and starts parallel execution.

#pragma omp parallel [*clause* [,] *clause* ...]
structured-block

```
#pragma omp parallel
{
    DoSomeWork (res, M) ;
    DoSomeOtherWork (res, M) ;
}
```

Pragma: for



for [2.7.1] [2.7.1]

Specifies that the iterations of associated loops will be executed in parallel by threads in the team in the context of their implicit tasks.

Pragma: for



- The **for** pragma can be used inside a block of code already marked with the **parallel** pragma
- Loop iterations should be divided among the active threads
- There is a barrier synchronization at the end of the **for** loop

```
#pragma omp parallel
{
    DoSomeWork(res, M);
    #pragma omp for
    for (i = 0; i < M; i++) {
        res[i] = huge();
    }
    DoSomeMoreWork(res, M);
}
```

Which loop to make parallel?



- Loop-carried dependence: dependence exists across iterations; i.e., if the loop is removed, the dependence no longer exists.
- Loop-independent dependence: dependence exists within an iteration; i.e., if the loop is removed, the dependence still exists.

Loop dependencies

```
for (i=1; i<n; i++) {  
    S1: a[i] = a[i-1] + 1;  
    S2: b[i] = a[i];  
}  
  
for (i=1; i<n; i++)  
    for (j=1; j< n; j++)  
        S3: a[i][j] = a[i][j-1] + 1;  
  
for (i=1; i<n; i++)  
    for (j=1; j< n; j++)  
        S4: a[i][j] = a[i-1][j] + 1;
```

$S1[i] \rightarrow_T S1[i+1]$: loop-carried

$S1[i] \rightarrow_T S2[i]$: loop-independent

$S3[i, j] \rightarrow_T S3[i, j+1]$:

- loop-carried on **for** j loop
- no loop-carried dependence in **for** i loop

$S4[i, j] \rightarrow_T S4[i+1, j]$:

- no loop-carried dependence in **for** j loop
- loop-carried on **for** i loop

Which Loop to Make Parallel?



```
main () {  
    int i, j, k;  
    float **a, **b;  
    ...
```

```
    for (k = 0; k < N; k++)
```

Loop-carried dependences

```
        for (i = 0; i < N; i++)
```

Can execute in parallel

```
            for (j = 0; j < N; j++)
```

Can execute in parallel

```
                a[i][j] = MIN(a[i][j], a[i][k] + a[k][j]);
```

Floyd's algorithm

Minimizing Threading Overhead



- There is a fork/join for every instance of
- Since fork/join is a source of overhead, we want to maximize the amount of work done for each fork/join
- Hence we choose to make the middle loop parallel
 - n fork/joins
 - For inner loop parallel, n^2 fork/joins

```
#pragma omp parallel for  
for ( ) {  
    ...  
}
```

Almost Right, but Not Quite

- `main () {`
- `int i, j, k;`
- `float **a, **b;` **Problem: j is a shared variable**
- `...`
- `for (k = 0; k < N; k++)`
- **`#pragma omp parallel for`**
- `for (i = 0; i < N; i++)`
- `for (j = 0; j < N; j++)`
- `a[i][j] = MIN(a[i][j], a[i][k] + a[k][j]);`

Problem Solved with private Clause



- `main () {`
 - `int i, j, k;`
 - `float **a, **b;`
 - `...`
 - `for (k = 0; k < N; k++)`
 - `#pragma omp parallel for private (j)`
 - `for (i = 0; i < N; i++)`
 - `for (j = 0; j < N; j++)`
 - `a[i][j] = MIN(a[i][j], a[i][k] + a[k][j]);`
- private(list)**
Declares one or more list items to be private to a task or a SIMD lane. Each task that references a list item that appears in a **private** clause in any statement in the construct receives a new list item.
- Tells compiler to make listed variables private

The Private Clause

- Reproduces the variable for each thread
 - Variables are un-initialized; C++ object is default constructed
 - Any value external to the parallel region is undefined

```
void work(float* c, int N)
{
    float x, y; int i;
    #pragma omp parallel for private(x,y)
    for(i = 0; i < N; i++) {
        x = a[i]; y = b[i];
        c[i] = x + y;
    }
}
```

Example: Dot Product



- Why won't the use of the **private** clause work in this example?

```
float dot_prod(float* a, float* b, int N)
{
    float sum = 0.0;
    #pragma omp parallel for private(sum)
    for(int i = 0; i < N; i++) {
        sum += a[i] * b[i];
    }
    return sum;
}
```

Reductions



- Given associative binary operator \oplus the expression
 - $a_1 \oplus a_2 \oplus a_3 \oplus \dots \oplus a_n$
- is called a *reduction*

OpenMP reduction Clause



- Reductions are so common that OpenMP provides a reduction clause for the `parallel for` pragma
- **reduction (op : list)**
- A PRIVATE copy of each list variable is created and initialized depending on the “op”
 - The identity value “op” (e.g., 0 for addition)
- These copies are updated locally by threads
- At end of construct, local copies are combined through “op” into a single value and combined with the value in the original SHARED variable

reduction(*reduction-identifier*:*list*)

Specifies a *reduction-identifier* and one or more list items. The *reduction-identifier* must match a previously declared *reduction-identifier* of the same name and type for each of the list items.

Operators for reduction (initialization values)			
+	(0)		(0)
*	(1)	^	(0)
-	(0)	&&	(1)
&	(~0)		(0)
max (Least representable number in reduction list item type)			
min (Largest representable number in reduction list item type)			

Reduction Example

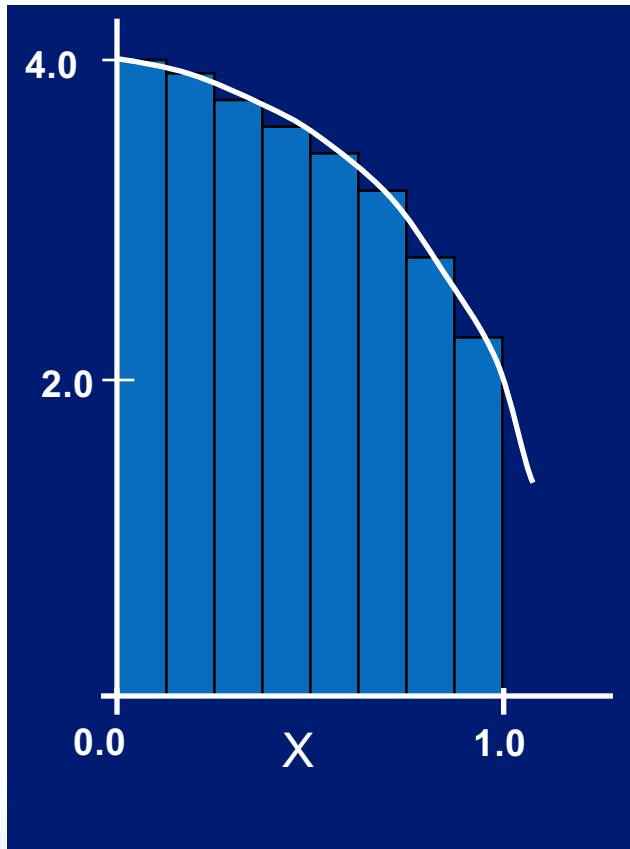


- Local copy of **sum** for each thread
- All local copies of sum added together and stored in shared copy

```
float dot_prod(float* a, float* b, int N)
{
    float sum = 0.0;
    #pragma omp parallel for reduction(+:sum)
    for(int i = 0; i < N; i++) {
        sum += a[i] * b[i];
    }
    return sum;
}
```

Numerical Integration Example

$$\int_0^1 \frac{4}{1+x^2} dx = \pi$$



```
static long num_rects=100000;  
double width, pi;
```

```
void main()
```

```
{  int i;  
    double x, sum = 0.0;
```

```
    width = 1.0/(double) num_rects;  
    for (i = 0; i < num_rects; i++){  
        x = (i+0.5)*width;  
        sum = sum + 4.0/(1.0 + x*x);  
    }
```

```
    pi = width * sum;  
    printf("Pi = %f\n",pi);
```

```
}
```

Numerical Integration: What's Shared?



```
static long num_rects=100000;  
double width, pi;
```

What variables can be shared?

```
void main()  
{  int i;  
    double x, sum = 0.0;                width, num_rects  
  
    width = 1.0/(double) num_rects;  
    for (i = 0; i < num_rects; i++){  
        x = (i+0.5)*width;  
        sum = sum + 4.0/(1.0 + x*x);  
    }  
    pi = step * sum;  
    printf("Pi = %f\n",pi);  
}
```

Numerical Integration: What's Private?



```
static long num_rects=100000;  
double width, pi;
```

What variables need to be private?

```
void main()  
{  int i;                                x, i  
    double x, sum = 0.0;  
  
    width = 1.0/(double) num_rects;  
    for (i = 0; i < num_rects; i++){  
        x = (i+0.5)*width;  
        sum = sum + 4.0/(1.0 + x*x);  
    }  
    pi = width * sum;  
    printf("Pi = %f\n",pi);  
}
```

Numerical Integration: Any Reductions?



```
static long num_rects=100000;  
double width, pi;
```

What variables should be set up for reduction?

```
void main()  
{  int i;  
    double x, sum = 0.0;           sum  
  
    width = 1.0/(double) num_rects;  
    for (i = 0; i < num_rects; i++){  
        x = (i+0.5)*width;  
        sum = sum + 4.0/(1.0 + x*x);  
    }  
    pi = step * sum;  
    printf("Pi = %f\n",pi);  
}
```

Solution to Computing Pi



```
static long num_rects=100000;
double width, pi;

void main()
{  int i;
   double x, sum = 0.0;
   #pragma omp parallel for private(x) reduction(+:sum)
   width = 1.0/(double) num_rects;
   for (i = 0; i < num_rects; i++){
       x = (i+0.5)*width;
       sum = sum + 4.0/(1.0 + x*x);
   }
   pi = step * sum;
   printf("Pi = %f\n",pi);
}
```