



Debug MPI Applications

Ruini Xue

School of Computer Science and Engineering

University of Electronic Science and Technology of China

2016

Outline



- Background
- Subgroup Reproducible Replay
- Challenges
- Evaluation
- Conclusion



Background

Difficulties of Debugging MPI



- Debugging parallel codes can be incredibly difficult, particularly as codes scale upwards.
 - Large number of processes
 - Distributed over different nodes
 - Long running time
 - Complicated communication pattern
 - All bugs in *sequential* programs
 - And new bugs due to *parallelism*
 - *Non-determinism, dead-lock...*
- **MPI debugging is a hard experience!**

- Static Checking
- Runtime Detection
- Postmortem Trace Analysis
- Model Checking
- Data Mining
- Replay-based Methods
 - Data-replay: Huge Log size
 - Order-replay: Huge replay resource
- ...
- **Debugging with replay is the future.** (The Parallel Computing Landscape: A view From Berkeley 2.0)

Static Checking



- Compile-time source checking
 - Limited to semantic errors
 - E.g., MPI-CHECK(CCPE'03)

- Cyclic debugging
 - ☹ Scalability, reproducibility, portability
 - E.g., *printf*, Totalview, DDT, PGDBG, Net-dbx
- PMPI Intercepting
 - ☹ only MPI APIs and predefined errors
 - E.g., Umpire(SC'00), MARMOT(Parco'03), Retrospect(EuroPVM/MPI'07)
- Debug version of MPI Library
 - ☹ Transparency, reproducibility
 - E.g., NEC mpi, MPICH,

Postmortem trace analysis



- Random message payload on non-deterministic point
 - E.g., HASE'05, IPDPS'07
- Check the trace for predefined errors
 - E.g., Intel Message Checker(SE-HPCS'05)

Replay-based Methods



- Data-driven replay
 - No code modification
 - PDT(TechRept'95):PMPI
 - Dieter(EuroPVM/MPI'01):MPICH-1.2.1
 - Retrospect (EuroPVM/MPI'07): OpenMPI core
 - MPI prototypes changed
 - Christian(PhDThesis'00): wrap MPI calls
- Order-driven replay
 - Happen-before relationship
 - NOPE (ACPC'99): only p2p APIs

- Model-Checking
 - Spin-MPI (VMCAI'07, EuroPVM/MPI'07)

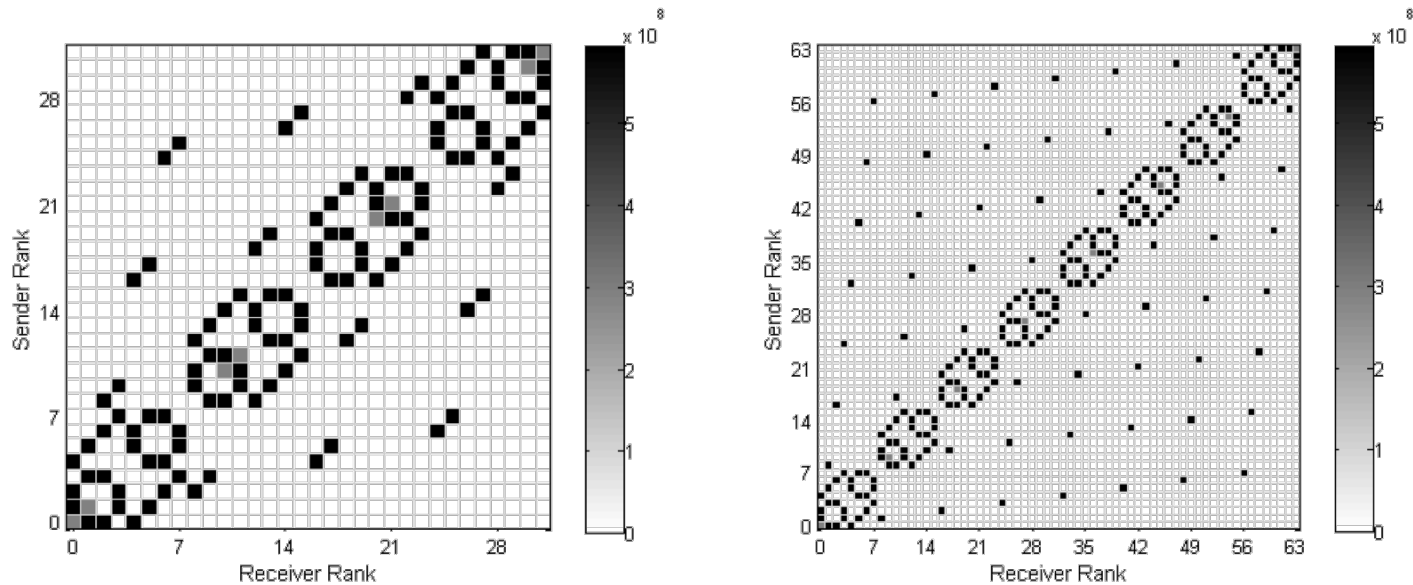


Subgroup Reproducible Replay

Inspiration: Communication Locality

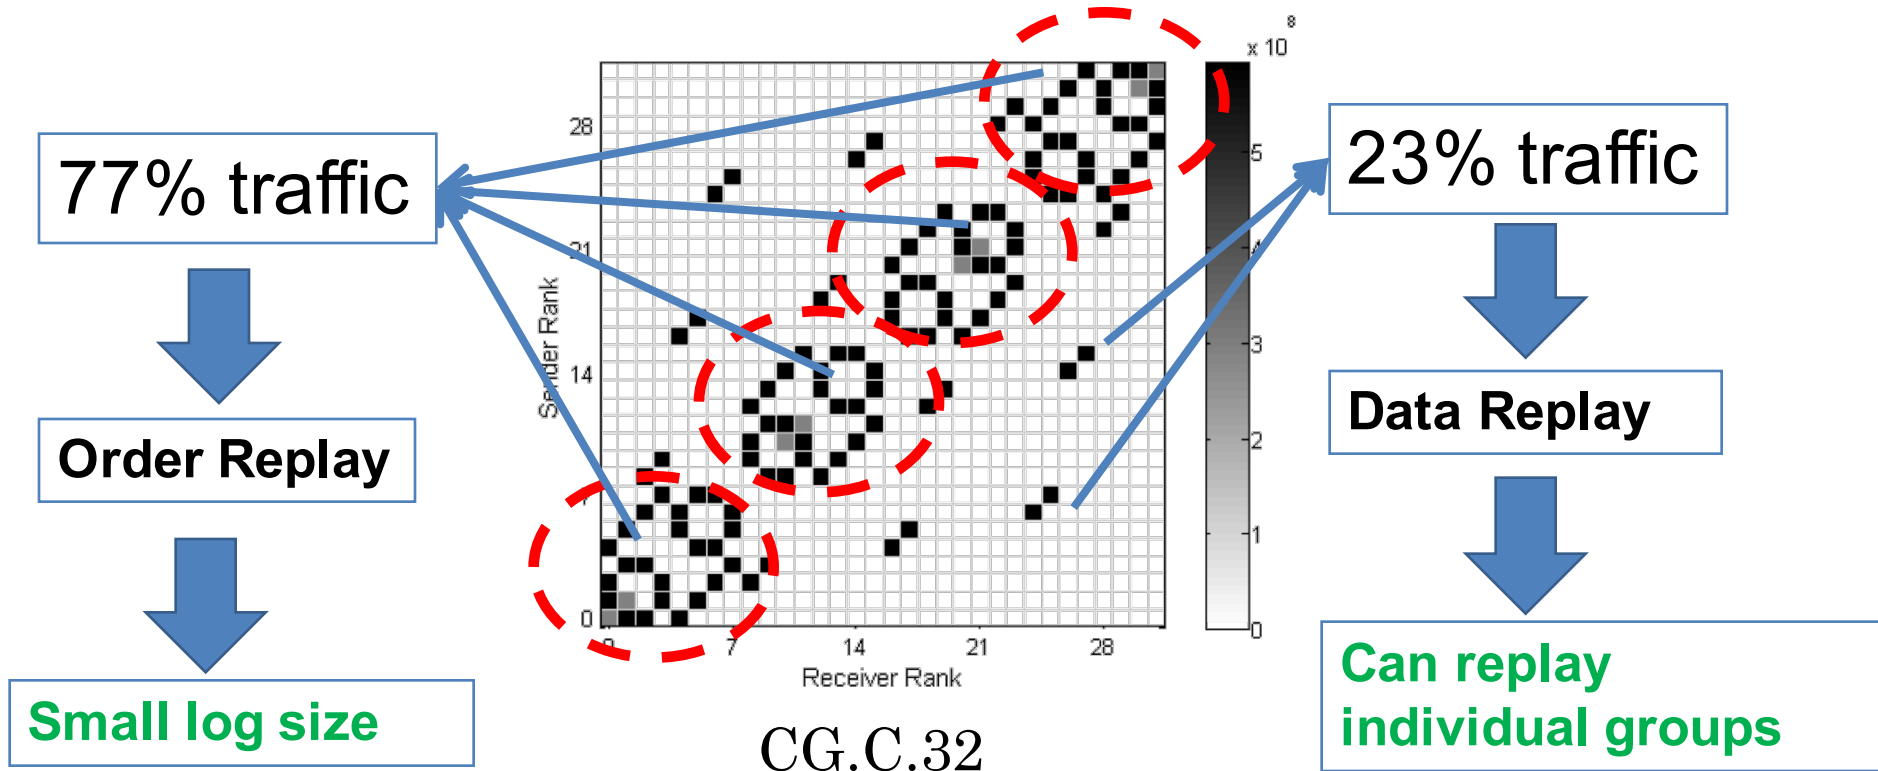


- NPB CG.C.32 & CG.C.64



- Intra-group communication: $\approx 77\%$

Observation: Communication Locality





Process grouping can be exploited for **a good balance** between record and replay cost (log size vs replay capability with limited resource)

Subgroup Reproducible Replay



- Combine data- and order-replay
 - Data-replay: inter-group communication
 - Order-replay: intra-group communication




- Generalization of them
 - A balance for record/replay trade-offs

Profile Run




- Collect communication trace (optional)
 - Profiling
- Expert Knowledge




Process Grouping



- Communication Graph Partition
 - Exploit comm. Locality
 - Process --- vertex
 - Comm. Volume --- edge weight





- Assign as your will



Record Phase


- Record different groups independently
- Log size decreases dramatically.



Replay



- Replay processes of one group altogether
- Feed-back
 - Inter-group comm. & OS calls
- Reproduce
 - Intra-group comm.
 - *faked* messages





So far, so easy 😊

Designs are cheap, but implementations are expensive.

—— *How (and How Not) to Write a Good Systems Paper, 1983*



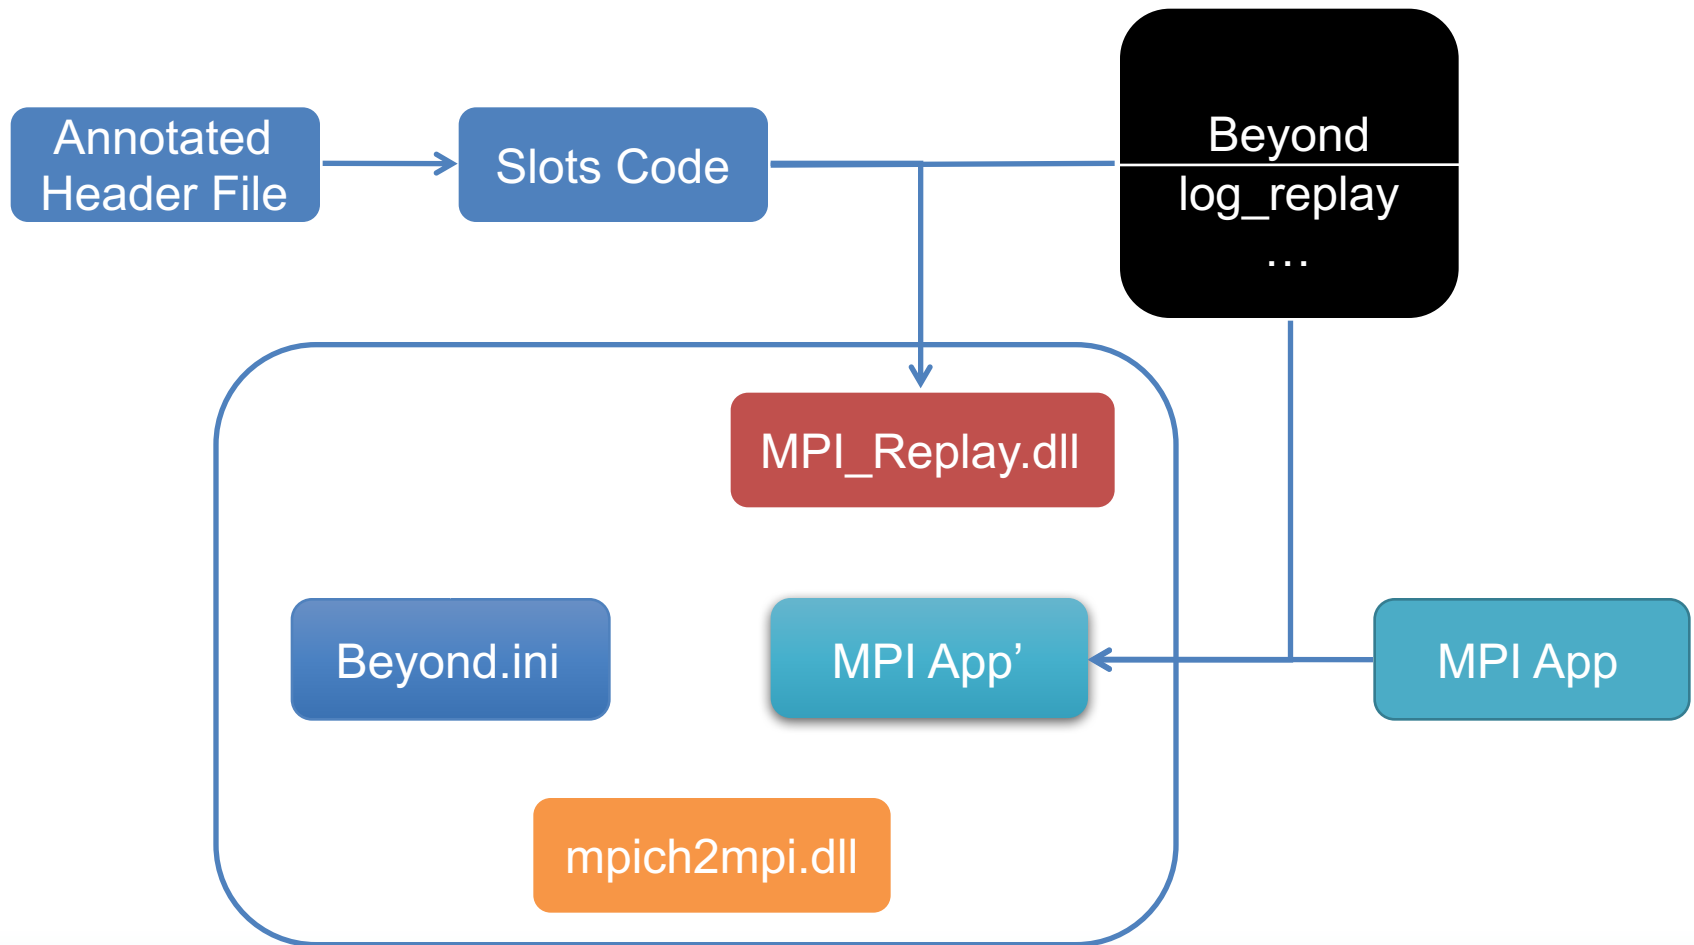
Challenges

Basic Challenges



- Interposition
 - Dynamic instrumentation
- Wrapper codes
 - Automatic generate = Annotation + Template
 - 18,000 C++ codes for MPI calls
- Faked replay
 - Replayer: a minimal MPI run-time
- MPI Standard
 - Dive in to details, too many tricks

How does it work?



Annotation



```
int
MPI_Init (
    [in] int *argc,
    [in] char ***argv
);
```

```
int
MPI_Get_processor_name (
    [out, ecap(*len), esize(*len+1), force] char *name,
    [out] int *len
);
```

```
int
MPI_Recv (
    [out, bsize("get_MPI_general_buf_size(datatype, count)", force] void* buf,
    [in] int count,
    [in] MPI_Datatype datatype,
    [in] int src,
    [in] int tag,
    [in] MPI_Comm comm,
    [out, opt(MPI_STATUS_IGNORE)] MPI_Status* status
);
```

Generated Record/Replay Codes (1)



```
/* int[cc(__cdecl), module(mpich2mpi)]
```

```
MPI_Init (
```

```
    [in] int * argc,
```

```
    [in] char * * * argv);*/
```

```
BEGIN_SIGEX_MIXIN(MPI_Init, hw_log_MPI_Init)
```

```
    START_PROFILE_WITH_NAME("hw_log_MPI_Init")
```

```
    g_LogProducer.set_identifier(g_papi_info_MPI_Init->api_sig);
```

```
    g_LogProducer << sx_last_ret;
```

```
    g_LogProducer << log::end;
```

```
END_SIGEX_MIXIN
```

```
BEGIN_SIGEX_SLOT(MPI_Init, hw_replay_MPI_Init)
```

```
    START_PROFILE_WITH_NAME("hw_replay_MPI_Init")
```

```
    g_LogConsumer.chk_identifier(g_papi_info_MPI_Init->api_sig);
```

```
    int sx_last_ret;
```

```
    g_LogConsumer >> sx_last_ret;
```

```
    g_LogConsumer >> log::end;
```

```
    return sx_last_ret;
```

```
END_SIGEX_SLOT
```


Generated Record/Replay Codes (2)



```
/*int[cc(__cdecl), module(mpich2mpi)]
MPI_Comm_size ( [in] MPI_Comm comm,   [out] int * size );*/

BEGIN_SIGEX_MIXIN(MPI_Comm_size, hw_log_MPI_Comm_size)
    START_PROFILE_WITH_NAME("hw_log_MPI_Comm_size")
    g_LogProducer.set_identifier(g_papi_info_MPI_Comm_size->api_sig);
    g_LogProducer << sx_last_ret;
    g_LogProducer << (uint32)sizeof(deref<int *>::type);
    g_LogProducer.write((LPVOID)size, (uint32)sizeof(deref<int *>::type));
    g_LogProducer << log::end;
END_SIGEX_MIXIN

BEGIN_SIGEX_SLOT(MPI_Comm_size, hw_replay_MPI_Comm_size)
    START_PROFILE_WITH_NAME("hw_replay_MPI_Comm_size")
    g_LogConsumer.chk_identifier(g_papi_info_MPI_Comm_size->api_sig);
    int sx_last_ret;
    g_LogConsumer >> sx_last_ret;
    uint32 nParamBufLength;
    g_LogConsumer >> nParamBufLength;
    g_LogConsumer.read((LPVOID)size, nParamBufLength);
    g_LogConsumer >> log::end;
    return sx_last_ret;
END_SIGEX_SLOT
```

More Challenges: Non-determinisms



- System calls
 - `gettimeofday`, `random`, `socket`...
 - data-replay works
- MPI calls
 - Inter-group Messages
 - Wildcard Receives
 - Waits, Test and Probes
 - Collective Operations
- How to record and replay them faithfully?

Inter-group Messages



- Are they non-deterministic?
- Record
 - Ignore **out** messages
 - Record **in** messages
 - include the order, if necessary
- Replay
 - Ignore **out** messages
 - Feed-back **in** messages from log

Wildcard Receives



- `MPI_ANY_SOURCE`, `MPI_ANY_TAG`
 - Record the real values, replace them during replay
- `MPI_STATUS[ES]_IGNORE`
 - Replace the parameter with a new allocated memory
 - Record the status data
 - Real values for `source` and `tag`
- Non-blocking Receives
 - Same solution but in `MPI_Wait/MPI_Test`

Waits, Tests, and Probes



- MPI_*some, *_*any (* = Wait | Test)
- Record
 - Map table: request <-> buffer info
 - Record returned request, and corresponding buffer
- Replay
 - Map table too
 - Feed-back (or receive) buffer
- Nothing special for Probes
 - following operation does the real work

- `MPI_BOTTOM`
 - Seldom used in real applications!
- `MPI_IN_PLACE`
 - Optimization for some collective operations (`MPI_Scatter`)
 - Just mark as: `opt(MPI_IN_PLACE)`
- `MPI_Start()`
 - Why no `MPI_Stop()`?
 - We can cleanup the rubbish, but not in a decent manner.
- `MPI_Alloc_mem()/MPI_Free_mem()`
 - The MPI forum committee is, eh..., dreaming.
- `MPI_Cancel()`
 - Actually as normal

1. What's the length of the buffer?



- `MPI_Recv(buf, count, MPI_FLOAT, src, tag, &req)`
 - `count * sizeof(MPI_FLOAT) =?` `count * sizeof(FLOAT)`

```
int MPI_Recv(void* buf, int count, MPI_Datatype datatype,  
             int source, int tag, MPI_Comm comm,  
             MPI_Status *status )
```

```
typedef int MPI_Datatype
```

- `sizeof()` sucks for even built-in/primitive datatypes!

Even worse



- Derived datatype
 - $\text{Typemap} = \{(\text{type}_0, \text{disp}_0), \dots, (\text{type}_{n-1}, \text{disp}_{n-1})\}$
 - $\text{Extent}(\text{Typemap}) = \text{ub}(\text{Typemap}) - \text{lb}(\text{Typemap})$
 - Size: remove the spaces in extent
 - Built-in datatypes are indeed derived datatypes!

```
#define MPI_INT ((MPI_Datatype)6)

MPI_INT = {(int, 0)}
```


Why extent and size?



- Insights on MPI_Send and MPI_Recv
 - $\text{sendbuf, recvbuf} = \text{count} * \text{extent}(\text{datatype}) + \text{lb}(\text{datatype})$
 - $\text{network data bytes} = \text{count} * \text{size}(\text{datatype})$
 - straightforward optimization
- *How do we know the size of a type?*
 - `int MPI_Type_size(MPI_Datatype dt, int* size)`

Determine the buffer size

- Solution
 - $\text{save count} * \text{size}(\text{datatype}) + \text{lb}(\text{datatype})$
 - usually: $\text{lb} = 0$

```
int
MPI_Recv (
    [out, bsize("get_MPI_buf_size(datatype, count)"), force] void* buf,
    .....,
    [out, opt(MPI_STATUS_IGNORE)] MPI_Status* status
);
```

2. What does the returned length mean?



- The length of the *'out'* buffer is indicated by another *'out'* parameter

```
int MPI_Get_processor_name (char *name, int *len);
```

- Old method:

```
int  
MPI_Get_processor_name (  
    [out, ecap(*len), esize(*len + 1)] char *name,  
    [out] int *len  
);
```

- **Error!**

Why is it wrong?



- log slot

```
BEGIN_SIGEX_MIXIN(MPI_Get_processor_name, hw_log_MPI_Get_processor_name)
  START_PROFILE_WITH_NAME("hw_log_MPI_Get_processor_name")
  g_LogProducer.set_identifier(g_papi_info_MPI_Get_processor_name->api_sig);
  g_LogProducer << sx_last_ret;
  if ((*len + 1)*sizeof(deref<char *>::type) > 0) {
    g_LogProducer << (uint32)(* len + 1)*sizeof(deref<char *>::type);
    g_LogProducer.write((LPVOID)name,
                        (uint32)(* len + 1)*sizeof(deref<char *>::type));
  }
  g_LogProducer << (uint32)sizeof(deref<int *>::type);
  g_LogProducer.write((LPVOID)len, (uint32)sizeof(deref<int *>::type));
  g_LogProducer << log::end;
END_SIGEX_MIXIN
```

Why is it wrong?



- replay slot

```
BEGIN_SIGEX_SLOT(MPI_Get_processor_name, hw_replay_MPI_Get_processor_name)
  START_PROFILE_WITH_NAME("hw_replay_MPI_Get_processor_name")
  g_LogConsumer.chk_identifier(g_papi_info_MPI_Get_processor_name->api_sig);
  int sx_last_ret;
  g_LogConsumer >> sx_last_ret;
  uint32 nParamBufLength;
  if ((*len + 1)*sizeof(deref<char *>::type) > 0) {
    g_LogConsumer >> nParamBufLength;
    g_LogConsumer.read((LPVOID)name, nParamBufLength);
  }
  g_LogConsumer >> nParamBufLength;
  g_LogConsumer.read((LPVOID)len, nParamBufLength);
  g_LogConsumer >> log::end;
  return sx_last_ret;
END_SIGEX_SLOT
```

Determine the buffer size



- Solution
 - a new tag: **force**
 - **always save it on log, and no test on replay**

```
int
MPI_Get_processor_name (
    [out, ecap(*len), esize(*len + 1), force] char *name,
    [out] int *len
);
```

New Log slot



```
BEGIN_SIGEX_MIXIN(MPI_Get_processor_name, hw_log_MPI_Get_processor_name)
  START_PROFILE_WITH_NAME("hw_log_MPI_Get_processor_name")
  g_LogProducer.set_identifier(g_papi_info_MPI_Get_processor_name->api_sig);
  g_LogProducer << sx_last_ret;
  g_LogProducer << (uint32)(* len + 1)*sizeof(deref<char *>::type);
  g_LogProducer.write((LPVOID)name,
                      (uint32)(* len + 1)*sizeof(deref<char *>::type));
  g_LogProducer << (uint32)sizeof(deref<int *>::type);
  g_LogProducer.write((LPVOID)len, (uint32)sizeof(deref<int *>::type));
  g_LogProducer << log::end;
END_SIGEX_MIXIN
```

New Replay slot



```
BEGIN_SIGEX_SLOT(MPI_Get_processor_name, hw_replay_MPI_Get_processor_name)
  START_PROFILE_WITH_NAME("hw_replay_MPI_Get_processor_name")
  g_LogConsumer.chk_identifier(g_papi_info_MPI_Get_processor_name->api_sig);
  int sx_last_ret;
  g_LogConsumer >> sx_last_ret;
  uint32 nParamBufLength;
  g_LogConsumer >> nParamBufLength;
  g_LogConsumer.read((LPVOID)name, nParamBufLength);
  g_LogConsumer >> nParamBufLength;
  g_LogConsumer.read((LPVOID)len, nParamBufLength);
  g_LogConsumer >> log::end;
  return sx_last_ret;
END_SIGEX_SLOT
```


3. What is expected for ‘opt’?

- *If the arg might be NULL, set its flag to ‘opt’*
- What about other special values but not NULL?
 - MPI_STATUS_IGNORE
 - MPI_STATUSES_IGNORE

```
MPI_Wait(&req, MPI_STATUS_IGNORE)
MPI_Waitall(cnt, &reqs, MPI_STATUSES_IGNORE)

#define MPI_STATUS_IGNORE (MPI_Status *)1    //!<= NULL
#define MPI_STATUSES_IGNORE (MPI_Status *)1 //!<= NULL
```

- *NULL is not enough!*

Extends opt



- Solution
 - provide test value for '*opt*'
 - [out, opt] arg
if (arg != NULL)
 - [out opt(MY_VALUE)] arg
if (arg != MY_VALUE)

```
int
MPI_Wait (
    [in] MPI_Request *request,
    [out, opt(MPI_STATUS_IGNORE)] MPI_Status *status
);
```

4. What do non-blocking operations imply?



- Return immediately, and test later before using
- Example

```
MPI_Isend(buf, cnt, dt, srg, tag, comm, &request);  
// do something  
MPI_Wait(&request, &status);
```

```
MPI_Irecv(buf, cnt, dt, srg, tag, comm, &request);  
// do something  
MPI_Wait(&request, &status);
```

- The buffer is **NOT** ready when it returns.

What do non-blocking operations imply?



- Similar to *asio*
- All non-blocking APIs post `MPI_Request` objects
- `MPI_Wait()` tests the request in blocking manner
 - `MPI_Wait/MPI_Waitall/MPI_Waitany/MPI_Waitsome`
- They set requests to `MPI_REQUEST_NULL` after returning
 - Take care of the pitfalls!
- *When and how to log the buffer?*
 - `MPI_Wait()` and `MPI_Test()` family APIs

- MPI opaque object
- Track the buffer info attached to the request
- Useful in receives only
 - non-blocking send is nothing different with blocking send
 - Harmless even if the send fails
- Should we know the request type?
 - Sure, MPI_Wait does NOT set it to MPI_REQUEST_NULL always. ([persistent] send/recv)

Additional Slots


- Solution

- Log

- save buffer info on MPI_Irecv() and MPI_Recv_init()



- save requests in MPI_Wait() family APIs
 - log the buffer after MPI_Wait() family APIs



- Replay

- back fill the buffer from the log as normal

5. What's special for collective operations?



- Who are they?
 - MPI_Bcast
 - MPI_Gather(v)/MPI_Allgather(v)
 - MPI_Scatter(v)
 - MPI_Reduce/MPI_Allreduce/MPI_Reduce_Scatter
 - MPI_Alltoall(v)
 - MPI_Scan
- *How to deal with non-blocking ones?*

What's special for collective operations?



- For the same arg, different meanings for different processes!
- e.g. MPI_Reduce()
 - **recvbuf** is only significant for the **root**!

```
int MPI_Reduce(void *sendbuf, void *recvbuf, int count,  
               MPI_Datatype datatype, MPI_Op op,  
               int root, MPI_Comm comm );
```

- Solution
 - always mark 'out', but different 'bsize'

What's special for collective operations?



- How to replay if some processes are out of the group?
 - Record phase: save the membership
 - Replay phase: emulated with p2p communication
 - MPI_Bcast() example

```
/* MPI_Bcast() replay code */
load MPI_Bcast rank_list from log
if (I am root) { /* for data sender */
    foreach rank in rank_list:
        if (rank is in replay group)
            send message to rank
} else { /* for data receiver */
    if (root is in group)
        recv message from root
    else
        load message from log
}
```

6. Where to pack the data?

- MPI_Pack/MPI_Unpack

```
int MPI_Pack(  
    [in] void *inbuf,  
    [in] int incount,  
    [in] MPI_Datatype datatype,  
    [out, bsize("get_pack_size()"), force] void *outbuf,  
    [in] int outcount,  
    [inout] int *position,  
    [in] MPI_Comm comm  
);
```

- Solution
 - Additional slots:



7. What about MPI_Cancel()?

- MPI_Cancel() just posts a cancel request!
- MPI_Wait() to make sure whether the request is cancelled or not internally by set the **cancelled** field in **status**.
- MPI_Test_cancelled() simply tests the **cancelled** field in **status**.

```
MPI_Cancel(&request);  
MPI_Wait(&request, &status);  
MPI_Test_cancelled(&status, &flag);
```

Remove the cancelled requests



- Solution
 - Filter the cancelled requests in MPI_Wait()/MPI_Test() family APIs when saving the non-blocking buffers according to the status.
 - Actually, this is an optimization more than a feature. Since there is no harm if we always save the cancelled operations.

8. Some other evils

- `MPI_BOTTOM`
 - No one would use this API in real applications
- `MPI_IN_PLACE`
 - Optimization for some collective operations (`MPI_Scatter`)
 - Just mark as: `opt(MPI_IN_PLACE)`
- `MPI_Start()`
 - Why no `MPI_Stop()`?
 - We can cleanup the rubbish, but not in a decent manner.
- `MPI_Alloc_mem()/MPI_Free_mem()`
 - The MPI forum committee is, eh..., dreaming.



9. I/O Redirection

- All STD IO HANDLES are redirected to socks
 - `printf` → `NtWriteFile`
- Replay
 - `printf` → `NtWriteConsoleA`
- Solution
 - Replayer helps
 - anonymous pipes to take over the STD HANDLES



Thanks!



Evaluation

Setup



- 8 nodes of 2×4 -Core Intel Xeon 2.33 GHz CPUs
- 8GB RAM
- 140GB Disk
- Windows Server 2003 Enterprise Edition SP1
- MPICH2-1.0.7
- 1Gbps Ethernet LAN

Applications




- Three kinds of communication patterns
- Different kinds of non-determinisms

Communication Pattern							
	Locality			All-to-All			M/S
Operations	CG	MG	LU	FT	GE	ASP	PU
Non-determ. MPI		√	√		√	√	√
Non-determ. Sys	√	√	√	√	√		
Coll. Operations	√	√	√	√	√	√	

Execution Time




- Record Overhead
- Replay Overhead



Log Size





- Locality: 38%
- All-to-All: 87%
- M/S: 98% (1%)



Group-size and Membership



- Informed vs. Random
- CG & FT (CLASS=C, NPROCS=64)





Thanks!