

# Transactional Memory: Programming Multi-Core Systems

Ruini Xue

School of Computer Science and Engineering
University of Electronic Science and Technology of China
2016

## A Trend in Hardware



### "The Movement to Multi-core Processors"

- Originates from inability to increase processor clock rate
- Profound impact on architecture, OS and applications

How to program multi-core systems effectively?

## Outline



- Overview of Multi-core
- Paradigms: Lock vs. Transaction
- Transactional Memory

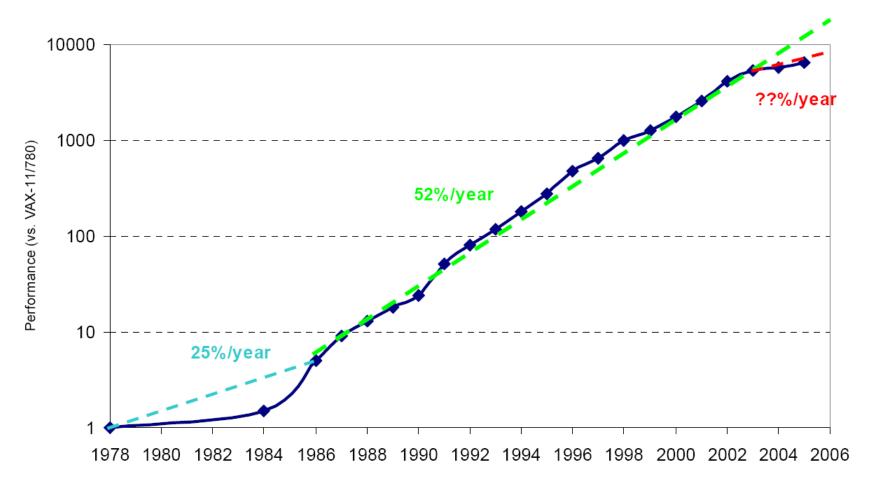
# Why Multi-Core?



- Areas of improving CPU performance in last 30 years
  - 1. Clock speed
  - 2. Execution optimization/ILP (Cycles-Per-Instruction)
  - 3. Cache
- All 3 are concurrency-agnostic
- Now
  - 1. Disappears: no Intel 4GHz CPU
  - 2. Slows down
  - 3. Still good

# **CPU Speed History**





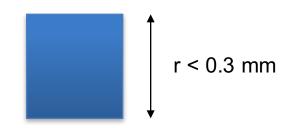
From "Computer Architecture: A Quantitative Approach", 4th edition, 2007

## Physical Distance



- Consider the 1 TFLOPS sequential machine:
  - Data must travel some distance, r, to get from memory to CPU.
  - To get 1 data element per cycle, this means  $10^{12}$  times per second at the speed of light,  $c = 3x10^8$  m/s. Thus  $r < c/10^{12}$  = 0.3 mm.

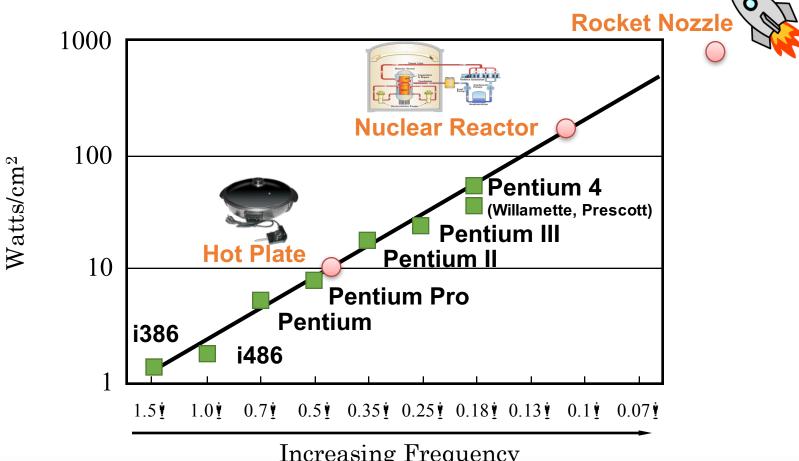
1 Tflop/s sequential machine



## Power Issue





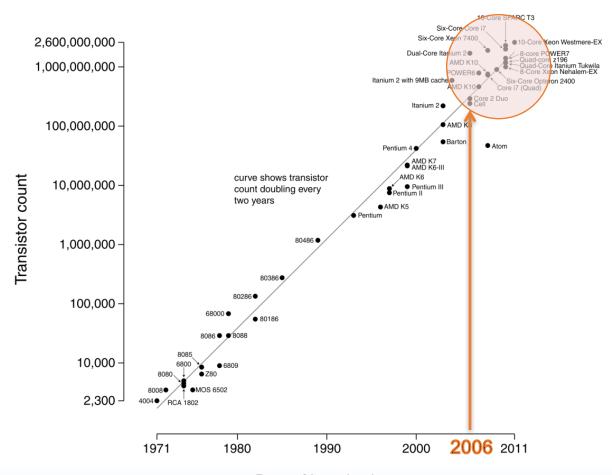


Increasing Frequency

## **Transistor Count**



#### Microprocessor Transistor Counts 1971-2011 & Moore's Law



#### Date of introduction

# Reality Today



- Multicore is the mainstream
- Exploit the performance
  - Let the cores run concurrently.
  - Design concurrent running units
    - multi-threads, multi-processes
- The software must be **concurrent**

"Concurrency is the next revolution in how we write software"
—— Hurb Sutter (2005): The Free Lunch is Over

# Costs/Problems of Concurrency



- pthread, MPI, PVM, openMP
- Overhead of locks, message passing
- Not all programs are parallelizable
- Programming concurrently is HARD
  - Complex concepts: mutex, read-write lock, queue...
  - Correct synchronization: race, deadlocks...
  - Getting speed-up

# Status in Synchronization



- Current mechanism: manual locking
  - Organization: lock for each shared structure
  - Usage: (block) → acquire → access → release
- Correctness issues
  - Under-locking → data races
  - Acquires in different orders → deadlock
- Performance issues
  - Difficult to find right granularity
  - Overhead of acquiring vs. allowed concurrency

## Transactions / Atomic Sections



- Databases has provided automatic concurrency control for 30 years: ACID transactions
- A transaction is a finite sequence of machine instructions executed by a single process, that satisfies the following properties
  - Atomicity
  - Isolation
  - Serialization only on conflicts
  - (optional) Rollback/abort support

UPDATE user set name="bar" where name="foo";

Question: Is it possible to provide database transaction semantics to general programming?

## How transactions work



- Prepare
  - Make private copy of shared data
- Work
  - Make updates on private copy
- Commit
  - If shared data is unchanged
    - Update shared data with private copy
  - Else conflict has occurred
    - Discard private copy and repeat transaction

## Transactions vs. Manual Locks



- Manual locking issues:
  - Under-locking
  - Acquires in different orders
  - Blocking
  - Conservative serialization (pessimistic)

- How transactions help:
  - No explicit locks
  - No ordering
  - Non-blocking, can cancel transactions
  - Serialization only on conflicts (optimistic)

## Transactions: simpler and more efficient

# Transactional Memory



- Attempts to simplify parallel programming by allowing a group of load and store instructions to execute in an atomic way.
- A concurrency control mechanism analogous to database transactions for controlling access to shared memory in concurrent computing.
- Simplifies concurrent programming significantly
  - Hardware TM
  - Software TM

## STM



- Transactions run in software
- A thread executes a transaction
- The transaction either commits or aborts
- Non-blocking The system makes progress
- Features
  - More flexible
  - Easier to modify and evolve
  - Integrate better with existing systems and languages

# Language Support



```
// Insert into a doubly-linked list atomically
atomic { // acquire lock
  newNode->prev = node;
  newNode->next = node->next;
  node->next->prev = newNode;
  node->next = newNode;
} // release lock
// Guard condition
atomic (queueSize > 0) { // conditional variable
  // remove item from queue and use it
```

# Direct Update STM



- Augment objects with (i) a lock, (ii) a version number
- Transactional write:
  - Lock objects before they are written to (abort if another thread has that lock)
  - Log the overwritten data we need it to restore the heap in case of retry, transaction abort, or a conflict with a concurrent thread
  - Make the update in place to the object
  - Increase the version numbers of object we've written, unlocking them

# Direct Update STM



- Transactional read:
  - Log the object's version number
  - Read from the object itself
  - · Check the version numbers of objects we've read
    - The same: commit
    - Different: abort



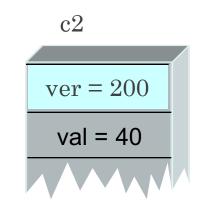
#### Thread T1

### int t = 0; atomic { t += c1.val; t += c2.val; }

#### Thread T2

```
atomic {
    t = c1.val;
    t ++;
    c1.val = t;
}
```

# c1 ver = 100 val = 10



T1's log:

T2's log:



```
Thread T1
```

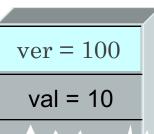
```
int t = 0;
atomic {
t += c1.val;
t += c2.val;
}
```

#### Thread T2

```
atomic {
    t = c1.val;
    t ++;
    c1.val = t;
}
```

T2's log:





c2

```
ver = 200
```

val = 40

```
T1's log:
```

c1.ver=100

T1 reads from c1: logs that it saw version 100



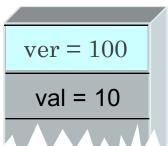
#### Thread T1

```
int t = 0;
atomic {
t += c1.val;
t += c2.val;
}
```

#### Thread T2

```
atomic {
    t = c1.val;
    t ++;
    c1.val = t;
}
```





c2

```
val = 40
```

ver = 200

T1's log:

c1.ver=100

T2's log:

c1.ver=100

T2 also reads from c1: logs that it saw version 100



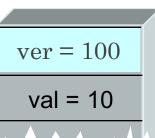
#### Thread T1

```
int t = 0;
atomic {
    t += c1.val;
    t += c2.val;
}
```

#### Thread T2

```
atomic {
    t = c1.val;
    t ++;
    c1.val = t;
}
```





c2

val = 40

ver = 200

```
T1's log:
```

c1.ver=100 c2.ver=200 T2's log:

c1.ver=100

Suppose T1 now reads from c2, sees it at version 200



#### Thread T1

```
int t = 0;
atomic {
    t += c1.val;
    t += c2.val;
}
```

#### Thread T2

```
atomic {
    t = c1.val;
    t ++;
    c1.val = t;
}
```

#### c1

locked:T2 val = 10

#### c2

ver = 200

val = 40

#### T1's log:

c1.ver=100 c2.ver=200

#### T2's log:

c1.ver=100 lock: c1, 100

Before updating c1, thread T2 must lock it: record old version number



#### Thread T1

```
int t = 0;
atomic {
    t += c1.val;
    t += c2.val;
}
```

#### Thread T2

```
atomic {
    t = c1.val;
    t ++;
    c1.val = t;
}
```

#### c1

locked:T2
val = 11

c2

ver = 200

val = 40

#### T1's log:

c1.ver=100 c2.ver=200

#### T2's log:

c1.ver=100 lock: c1, 100 c1.val=10 (2) After logging the old value, T2 makes its update in place to c1

(1) Before updating c1.val, thread T2 must log the data it's going to overwrite

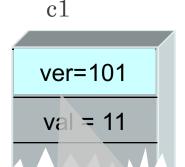


#### Thread T1

```
int t = 0;
atomic {
    t += c1.val;
    t += c2.val;
}
```

#### Thread T2

```
atomic {
    t = c1.val;
    t ++;
    c1.val = t;
}
```



c2

```
ver = 200
val = 40
```

#### T1's log:

c1.ver=100 c2.ver=200

#### T2's log:

c1.ver=100 lock: c1, 100 c1.val=10 (2) T2's transaction commits successfully. Unlock the object, installing the new version number

(1) Check the version we locked matches the version we previously read



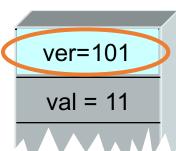
#### Thread T1

```
int t = 0;
atomic {
    t += c1.val;
    t += c2.val;
}
```

#### Thread T2

```
atomic {
    t = c1.val;
    t ++;
    c1.val = t;
}
```

#### c1



c2

ver = 200

val = 40

#### T1's log:

T2's log:

- (1) T1 attempts to commit. Check the versions it read are still up-to-date.
- (2) Object c1 was updated from version 100 to 101, so T1's transaction is aborted and re-run.

## State of the art



#### Hardware:

- Rock processor (canceled by Oracle)
- Blue Gene/Q processor from IBM (Sequoia supercomputer)<sup>[9]</sup>
- IBM zEnterprise EC12, the first commercial server to include transactional memory processor instructions
- Intel's Transactional Synchronization Extensions (TSX), available in select Haswell-based processors and newer
- IBM POWER8<sup>[10][11]</sup>

#### • Software:

- Vega 2 from Azul Systems<sup>[12]</sup>
- STM Monad in the Glasgow Haskell Compiler
- STMX in Common Lisp<sup>[13]</sup>
- Refs in Clojure
- gcc 4.7+ for C/C++<sup>[14][15][16][17]</sup>
- PyPy<sup>[18]</sup>

# Summary



- Why is multi-core architecture chosen?
- How does transaction simplify concurrency control?
- What is TM and how does STM work?

## Discussion



• Why TM is not popular in single-core era?

- STM challenges
  - Communications, or side-effects
    - File I/O, network......
  - Interrupts
  - Long transactions
  - Nested transactions



# **THANKS**



# Backup slides

# Compiler integration



- We expose decomposed log-writing operations in the compiler's internal intermediate code (no change to MSIL)
  - OpenForRead before the first time we read from an object (e.g. c1 or c2 in the examples)
  - OpenForUpdate before the first time we update an object
  - LogOldValue before the first time we write to a given field

```
Source code atomic {

...

t += n.value;
n = n.next;

...

Basic intermediate code code

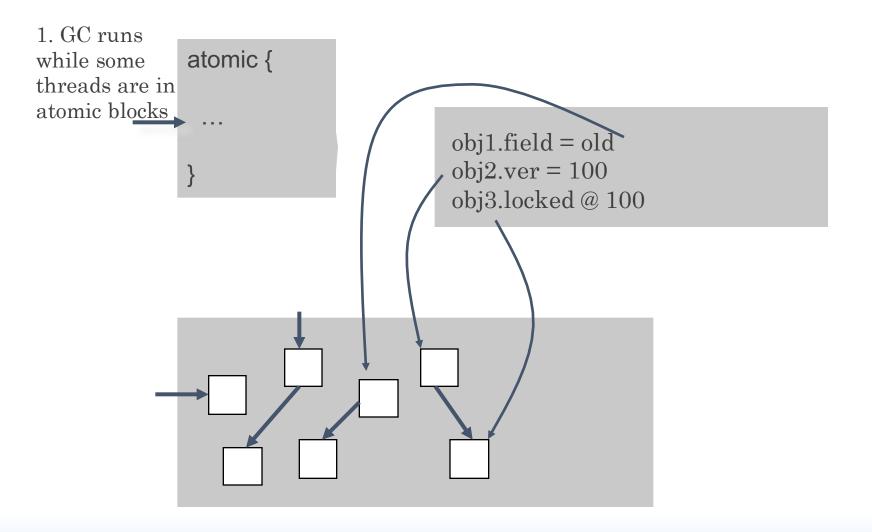
Optimized intermediate code code

OpenForRead(n);
t = n.value;
OpenForRead(n);
t = n.value;
n = n.next;

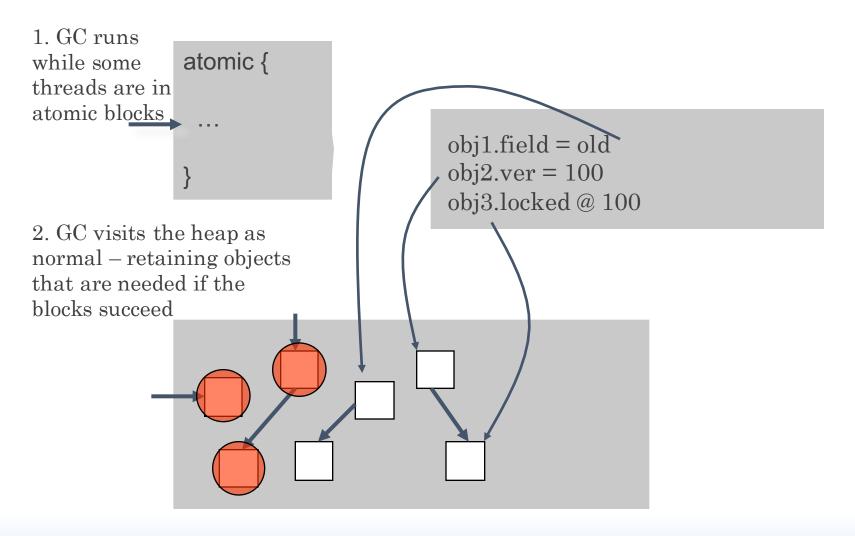
OpenForRead(n);
n = n.next;
```

# Runtime integration – garbage collection

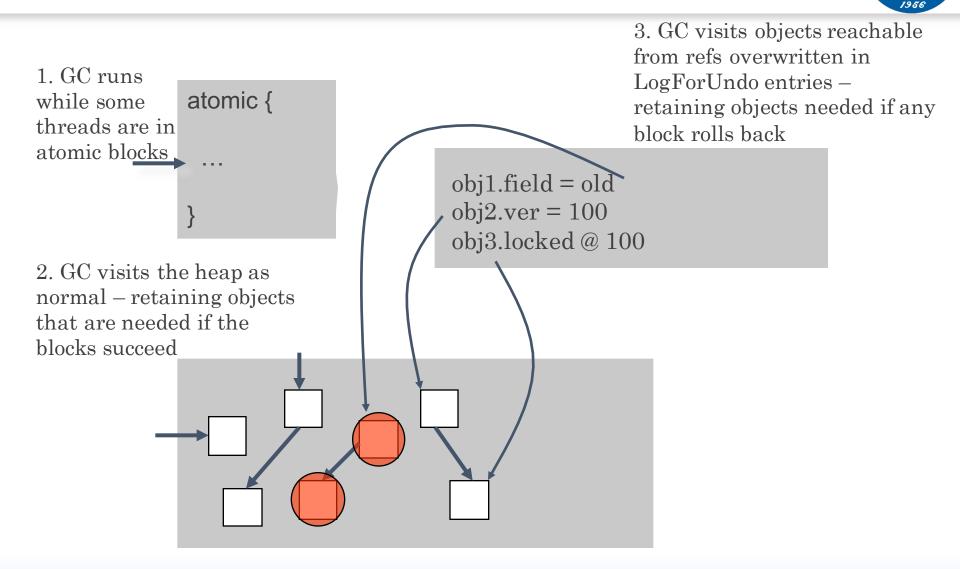




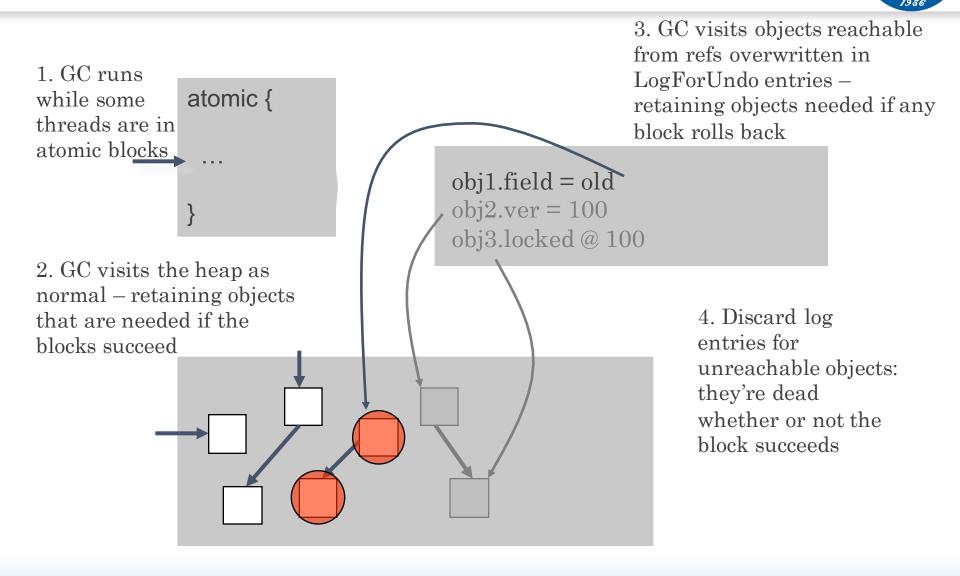




## Runtime integration – garbage collection

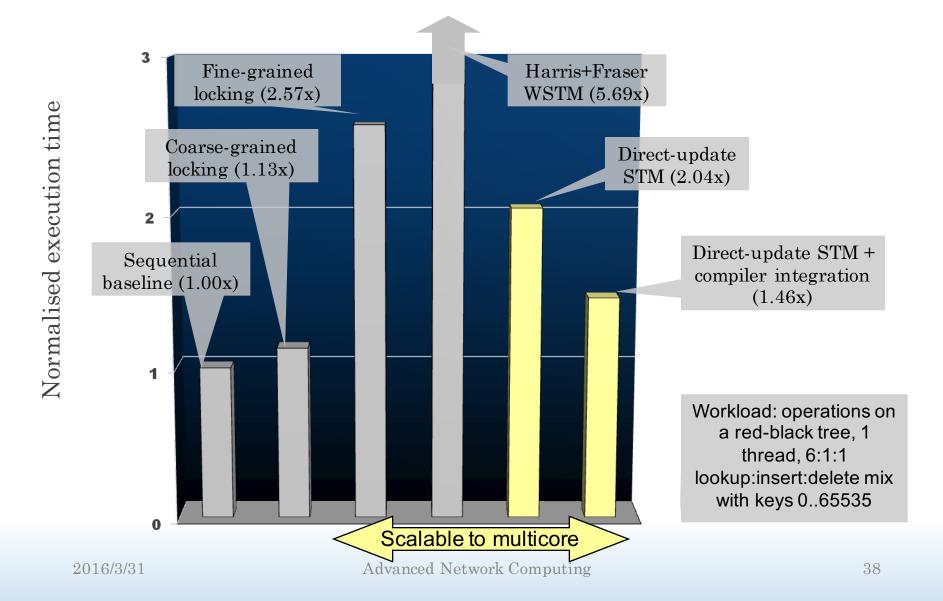


# Runtime integration – garbage collection



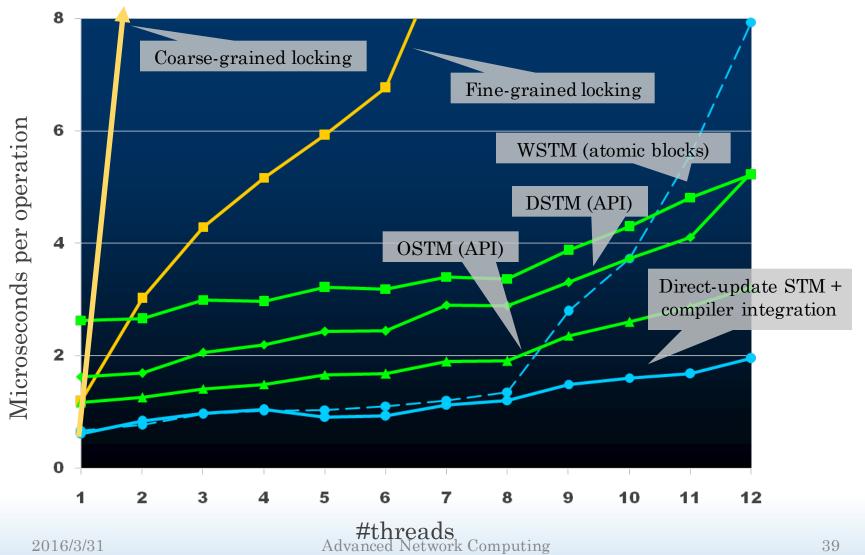
# Results: Against Previous Work





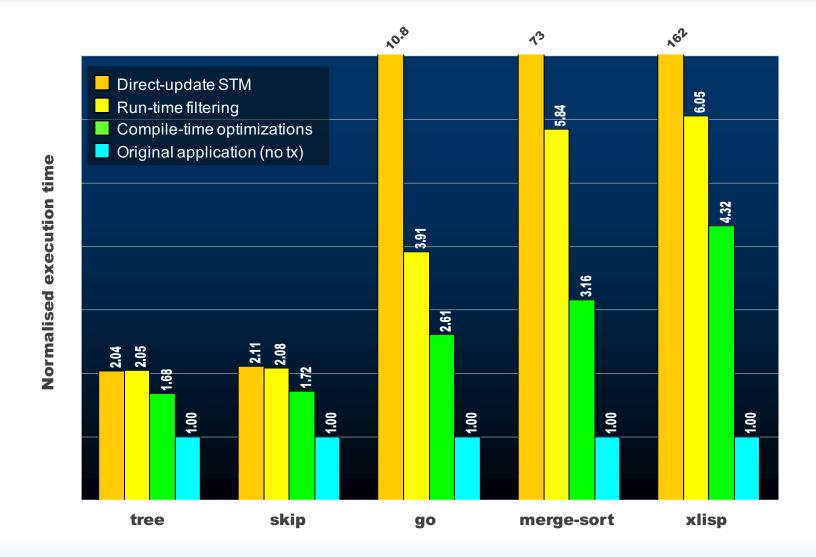
# Scalability (µ-benchmark)





# Results: long running tests





# Summary



- A pure software implementation can perform well and scale to vast transactions
  - Direct update
  - Pessimistic & optimistic CC
  - Compiler support & optimizations
- Still need a better understanding of realistic workload distributions

# Design Space



- Hardware Transactional Memory vs. software TM
- Granularity: object, word, block
- Update method
  - Deferred: discard private copy on aborts
  - Direct: control access to data, erase update on aborts
- Concurrency control
  - Pessimistic: prevent conflicts by locking
  - Optimistic: assumes no conflict and retry if there is

• ...

# Potential Multi-Core Apps



Application Category	Examples
Server apps w/o shared state	Apache web server
Server apps with shared state	MMORPG game server
Stream-Sort data processing	MapReduce, Yahoo Pig
Scientific computing (many different models)	BLAS, Monte Carlo, N-Body
Machine learning	HMM, EM algorithm
Graphics and games	NVIDIA Cg, GPU computing